

CS 6240- Large-scale Parallel Data Processing  
Homework 3  
Name: Srikanth Babu Mandru

**Words used in document:**

user1 – user at first position in edge

user2 - user at second position in edge

PATH2 – number of paths of length 2

TRIANGLES– number of TRIANGLES in twitter edges dataset

**Twitter Followers Count (Spark/Scala) Implementation:****Pseudo Codes:**

As the Spark/Scala code is very concise, I have explained the pseudo code in terms of Spark/Scala code (without “.” chaining”) with comments followed by for each statement. All the functions are applied in an order as mentioned in the pseudo-code.

**RDD-G:**

textFile ← Read the input file

textFile.

flatMap(line => line.split("\n")) // split the input records by “\n”

flatMap(line => line.split(",")(1)) //parse the records by using “,” separation &emit user2

map(word => (word, 1)) // assign “1” to each user

groupByKey() // group the records using “user2” as key

map(token => (token.\_1,token.\_2.sum)) //for each user, emit user id and count of followers

**RDD-R:**

textFile.flatMap (line => line.split("\n")) // split the input records by “\n”

map (line => (line.split(",")(1), 1 )) // parse records by using “,” separation  
// & emit (user2, count of “1”)

reduceByKey((a, b) => a+b) // For each user, sum the followers and emit the  
// (user id, followers count)

**RDD-F:**

textFile.

```

    flatMap(line => line.split("\n"))    // split the input records by "\n"

    map (line => (line.split(",")(1), 1 )) // parse records by using "," separation
                                           // & emit ( user2, count of "1")

    foldByKey(0)((a, b) => a + b )       // For each user, sum the followers starting from "0"
                                           // and emit the (user id, followers count)

```

**RDD-A:**

```

textFile.flatMap(line => line.split("\n"))    // split the input records by "\n"

    flatMap (line => line.split(",")(1))       //parse records by using "," separation & emit user2

    map (word => (word, 1))                    // assign "1" to each user

    aggregateByKey((k,v) => k+v , (k,v) => k+v) // For each user, sum the followers
                                           // and emit the (user id, followers count)

```

**DSET:**

input\_dataset ← read the input file into dataset with “user1”, “user2” as columns

```

input_dataset.groupBy($"user2").count()    // group all the records for each user2 and count
                                           //number of records which will be followers count

```

**Debug String Analysis:****RDD-G:**

```

(Debug info :,(10) CoalescedRDD[7] at coalesce at FollowerCount_RDDG.scala:29 []
| MapPartitionsRDD[6] at map at FollowerCount_RDDG.scala:29 []
| ShuffledRDD[5] at groupByKey at FollowerCount_RDDG.scala:28 []
+-(40) MapPartitionsRDD[4] at map at FollowerCount_RDDG.scala:27 []
| MapPartitionsRDD[3] at flatMap at FollowerCount_RDDG.scala:26 []
| MapPartitionsRDD[2] at flatMap at FollowerCount_RDDG.scala:26 []
| input MapPartitionsRDD[1] at textFile at FollowerCount_RDDG.scala:25 []
| input HadoopRDD[0] at textFile at FollowerCount_RDDG.scala:25 [])

```

**RDD-R:**

```
(Debug info :,(40) ShuffledRDD[4] at reduceByKey at FollowerCount_RDDR.scala:23 []
+- (40) MapPartitionsRDD[3] at map at FollowerCount_RDDR.scala:23 []
  | MapPartitionsRDD[2] at flatMap at FollowerCount_RDDR.scala:22 []
  | input MapPartitionsRDD[1] at textFile at FollowerCount_RDDR.scala:21 []
  | input HadoopRDD[0] at textFile at FollowerCount_RDDR.scala:21 [])
```

**RDD-F:**

```
(Debug info :,(40) ShuffledRDD[4] at foldByKey at FollowerCount_RDDF.scala:23 []
+- (40) MapPartitionsRDD[3] at map at FollowerCount_RDDF.scala:22 []
  | MapPartitionsRDD[2] at flatMap at FollowerCount_RDDF.scala:22 []
  | input MapPartitionsRDD[1] at textFile at FollowerCount_RDDF.scala:21 []
  | input HadoopRDD[0] at textFile at FollowerCount_RDDF.scala:21 [])
```

**RDD-A:**

```
(Debug info :,(40) ShuffledRDD[5] at aggregateByKey at FollowerCount_RDDA.scala:24 []
+- (40) MapPartitionsRDD[4] at map at FollowerCount_RDDA.scala:23 []
  | MapPartitionsRDD[3] at flatMap at FollowerCount_RDDA.scala:22 []
  | MapPartitionsRDD[2] at flatMap at FollowerCount_RDDA.scala:22 []
  | input MapPartitionsRDD[1] at textFile at FollowerCount_RDDA.scala:21 []
  | input HadoopRDD[0] at textFile at FollowerCount_RDDA.scala:21 [])
```

**Explain() Analysis for DSET:****Logical and Physical plan:**

== Parsed Logical Plan ==

```
'Aggregate [users2], [unresolvedalias('users2, None), count(1) AS count#24L]
+- Project [_c0#10 AS users1#14, _c1#11 AS users2#15]
  +- Relation[_c0#10,_c1#11] csv
```

== Analyzed Logical Plan ==

```
users2: string, count: bigint
Aggregate [users2#15], [users2#15, count(1) AS count#24L]
+- Project [_c0#10 AS users1#14, _c1#11 AS users2#15]
  +- Relation[_c0#10,_c1#11] csv
```

== Optimized Logical Plan ==

```
Aggregate [users2#15], [users2#15, count(1) AS count#24L]
+- Project [_c1#11 AS users2#15]
  +- Relation[_c0#10,_c1#11] csv
```

== Physical Plan ==

```

*(2) HashAggregate(keys=[users2#15], functions=[count(1)], output=[users2#15, count#24L])
+- Exchange hashpartitioning(users2#15, 200)
   +- *(1) HashAggregate(keys=[users2#15], functions=[partial_count(1)], output=[users2#15, count#29L])
      +- *(1) Project [_c1#11 AS users2#15]
         +- *(1) FileScan csv [_c1#11] Batched: false, Format: CSV, Location:
InMemoryFileIndex[file:/Users/srikanthmandru/my-system/semester
3/lsp/homeworks/HW3/Spark-traingl..., PartitionFilters: [], PushedFilters: [], ReadSchema:
struct<_c1:string>

```

### **Important points to note from analysis:**

From the above, it can be inferred that RDD-F(foldByKey), RDD-A (aggregatebykey), RDD-R (reducebykey) are similar to MapReduce's in-mapper combining.

### **(Twitter Social Amplifier)**

### **Triangles Count through Joins (Spark/Scala) Implementation:**

#### **(1) Reduce-side Join RDD Implementation:**

##### **Pseudo-code:**

input\_data ← read the data from input file

triangle\_counter ← 0

filterdata ← input\_data.filter( user1 < max\_filter and user2 < max\_filter ) // filter the input data to include

edges ← filterdata.map(user1 , user2) // only users with user ids less than max\_filter  
// emit tuples of (user1 , user2)

edges\_reversed ← filterdata.map(user2 , user1) // emit tuples of (user2 , user1)

path2\_edges ← edges\_reversed.join(edges) // Join two data with one table's key as "user2" and other  
// table's key as "user1". This forms PATH2 edges.

path3\_edges ← path2\_edges.join(edges\_reversed) // Join path2 edges with edges table in reverse order.  
// This forms Path3 edges

path3\_edges.foreach(users => if (user1 == user2){ triangle\_counter += 1 } ) // Check if users are same,  
// if same increment the triangle count by "1"

return ( triangle\_counter /3) // divide by "3" to get exact triangles

**(2) Reduce-side Join Dataset/ DataFrame Implementation:****Pseudo-code:**

```

twitterDS1 ← load (input file)
               .where(user1 < max_filter and user2 < max_filter) // Read input file into dataset and remove
                                                                    // edges with user ids greater than max_filter

twitterDS2 ← twitterDS1

PATH2DS ←
  twitterDS1.joinWith(twitterDS2, twitterDS1.user2 == twitterDS2.user1) // Join both the datasets with
                                                                    // keys as “user2” and “user1” of different datasets

TRIANGLEDS ←
  PATH2DS
  .joinWith(twitterDS1, PATH2DS.user1 == twitterDS1.user2 && PATH2DS.user2 == twitterDS1.user1 )
                                                                    // Join edges of length 2 with original edges and check if users are same
                                                                    // If same, keep the records in dataset.
return (size (TRIANGLEDS) / 3) // count the records in dataset and divide by “3” to get exact triangles

```

**(3) Replicated (Map-side) Join RDD Implementation:****Pseudo-code:**

```

input_data ← read the data from input file

triangle_counter ← 0

filterdata ← input_data.filter( user1 < max_filter and user2 < max_filter ) // filter the input data
to include                                                                    // only users with user ids less than max_filter

edges ← filterdata.map(user1 , user2)                                         // emit tuples of (user1 , user2)

edges_map ← broadcast(edges.groupby(user1)) //create Hashmap for each user1 and
                                                                    // broadcast Hashmap to each partition

// Now, Join the edges with HashMap “edges_map” two times. While joining recursively, check if first
// and last users are same and if same, increment the number of “TRIANGLES”
edges.map{ (user1 , user2) =>
  user2_set ← edges_map.get(user2)
  for each user3 ∈ user2_set:
    Increment “PATH2_counter” by 1
    user3_set ← edges_map.get(user3)
    for each user4 ∈ user3_set:
      if (user4 == user1)
        Increment “triangle_counter” by 1
}
return ( triangle_counter /3) // divide by “3” to get exact triangles

```

#### (4) Replicated (Map-side) Join Dataset/DataFrame Implementation:

##### Pseudo-code:

```

twitterDS1 ← load (input file)
               .where(user1 < max_filter and user2 < max_filter) // Read input file into dataset and remove
                                                                    // edges with user ids greater than max_filter

twitterDS2 ← broadcast(twitterDS1)                                // broadcast the twitter edges data over partitions

PATH2DS ←
  twitterDS1.join(twitterDS2, twitterDS1.user2 == twitterDS2.user1) // Join both the datasets with
                                                                    // keys as “user2” and “user1” of different datasets

TRIANGLEDS ←
  PATH2DS
  .joinWith(twitterDS1, PATH2DS.user1 == twitterDS1.user2 && PATH2DS.user2 == twitterDS1.user1 )
                                                                    // Join edges of length 2 with original edges and check if users are same
                                                                    // If same, keep the records in dataset.
return (size (TRIANGLEDS) / 3) // count the records in dataset and divide by “3” to get exact triangles

** Note: ‘broadcast(twitterDS1)’ makes spark to automatically choose the “partition + broadcast” join
type while using “join”.

```

##### Output:

Output can be inferred from the both syslog files provided in the “log files path” directory mentioned below and output files from the spark/scala program:

Configuration	Small Cluster Result (1 master, 5 Workers)	Large Cluster Result (1 master, 8 Workers)
<b>RS-R, MAX = 10000</b>	Running time: 4.87 minutes, Triangle count: 520296	Running time: 3.9 minutes, Triangle count: 520296
<b>RS-D, MAX = 10000</b>	Running time: 10.4 minutes, Triangle count: 520296	Running time: 9.13 minutes, Triangle count: 520296
<b>Rep-R, MAX = 10000</b>	Running time: 2.2 minutes, Triangle count: 520296	Running time: 2.17 minutes, Triangle count: 520296
<b>Rep-D, MAX = 10000</b>	Running time: 10.9 minutes, Triangle count: 520296	Running time: 8.73 minutes, Triangle count: 520296

**Note:** Here, for large cluster analysis, I couldn't able to get more than 8 workers. All programs are executed with similar setup as HW2 like machine instance type, max\_filter value, etc.

**Note:** For RDD-G, I got Java Heap error even after many attempts to solve the error. I have included the syslog for the run. Except that, all other programs ran successfully.

**Log files path:**

I have provided log files in the following directory separately for "Reduce side Join" and "Replicated join". Also, separated each with respect to "RDD" and "Dataset".

**Special mention:** Along with syslog from executions "steps" of aws, included syslog from aws containers as "\*" from container" beside respective syslog, where you can see the results of number of triangles and PATH2 by searching for "root" in the file.

HW3/logs/

**Output files path:**

HW3/output/ ; where '#' represents the run number

**Readme path:** (for execution steps or procedure)

HW3/Readme.txt

**Report path:**

HW3/Srikanth\_Mandru\_HW3.pdf